

Image Classification based on Convolutional Neural Network

Presented by:

Veronica Naosekpam

PhD Scholar


Computer Science & Engineering

Indian Institute of Information Technology Guwahati, Assam, India



Introduction

- Assigning an input image, one label from a fixed set of categories.
- Classify different types of sports (here, Soccer and Rugby) using convolution neural networks in two ways :
 - Training from scratch.
 - Transfer learning.
- Use Visual Geometry Group- VGG16- model.
 - A pre-trained model on ImageNet



Setting up custom Image Dataset

- Two popular sources : **ImageNet** and **Google OpenImages** via python scripts.
- Input : 3058
 - Train : 2448
 - Rugby : 1224
 - Soccer : 1224
 - Test: 610
 - Rugby : 305
 - Soccer : 305



Sample images from Rugby class



Sample images from Soccer class

- Import the required libraries.
- Matplotlib and Seaborn for visualizing our dataset
- Image data : OpenCV.

```
[1]: 1 import matplotlib.pyplot as plt
      2 import seaborn as sns
      3 import keras
      4 from keras.models import Sequential
      5 from keras.layers import Dense, Conv2D , MaxPool2D , Flatten , Dropout
      6 from keras.preprocessing.image import ImageDataGenerator
      7 from keras.optimizers import Adam
      8 from sklearn.metrics import classification_report, confusion_matrix
      9 import tensorflow as tf
     10 import cv2
     11 import os
     12 import numpy as np
     13
     14 os.environ["CUDA_DEVICE_ORDER"]="PCI_BUS_ID"
     15 os.environ["CUDA_VISIBLE_DEVICES"]="1"
```

Using TensorFlow backend.

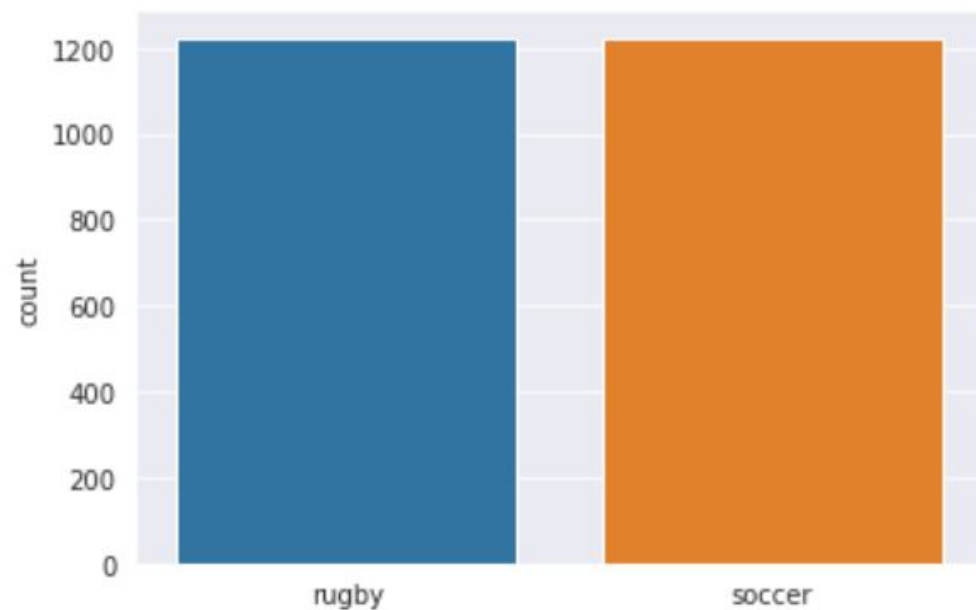
- Loading the data

```
1 labels = ['rugby', 'soccer']
2 img_size = 224
3 def get_data(data_dir):
4     data = []
5     for label in labels:
6         path = os.path.join(data_dir, label)
7         class_num = labels.index(label)
8         for img in os.listdir(path):
9             try:
10                 img_arr = cv2.imread(os.path.join(path, img))[...,:-1] #convert BGR to RGB format
11                 resized_arr = cv2.resize(img_arr, (img_size, img_size)) # Reshaping images to preferred size
12                 data.append([resized_arr, class_num])
13             except Exception as e:
14                 print(e)
15     return np.array(data)
```

```
1
2 train = get_data('/home/rs/veronica.naosekpam/Dataset/train')
3 val = get_data('/home/rs/veronica.naosekpam/Dataset/test')
4
```

```
1 l = []
2 for i in train:
3     if(i[1] == 0):
4         l.append("rugby")
5     else:
6         l.append("soccer")
7 sns.set_style('darkgrid')
8 sns.countplot(l)
```


5]: <AxesSubplot:ylabel='count'>



- Visualizing random images

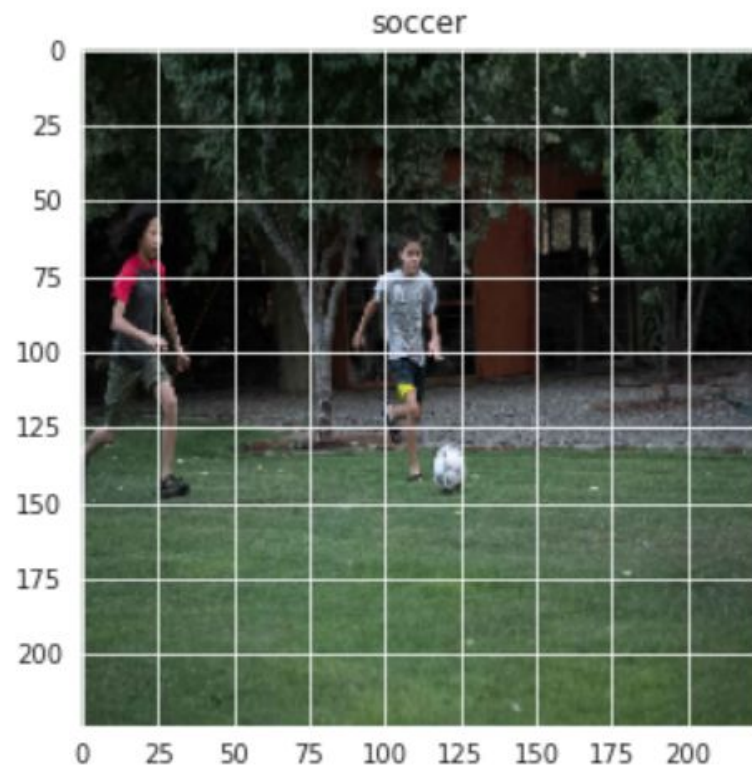
```
1 plt.figure(figsize = (5,5))
2 plt.imshow(train[1][0])
3 plt.title(labels[train[0][1]])
```

```
Text(0.5, 1.0, 'rugby')
```



```
1
2 plt.figure(figsize = (5,5))
3 plt.imshow(train[-4][0])
4 plt.title(labels[train[-1][1]])
```

```
Text(0.5, 1.0, 'soccer')
```



- Data Preprocessing and Data Augmentation.

```
1 x_train = []
2 y_train = []
3 x_val = []
4 y_val = []
5
6 for feature, label in train:
7     x_train.append(feature)
8     y_train.append(label)
9
10 for feature, label in val:
11     x_val.append(feature)
12     y_val.append(label)
13
14 # Normalize the data
15 x_train = np.array(x_train) / 255
16 x_val = np.array(x_val) / 255
17
18 x_train.reshape(-1, img_size, img_size, 1)
19 y_train = np.array(y_train)
20
21 x_val.reshape(-1, img_size, img_size, 1)
22 y_val = np.array(y_val)
```



```
1 datagen = ImageDataGenerator(  
2     featurewise_center=False, # set input mean to 0 over the dataset  
3     samplewise_center=False, # set each sample mean to 0  
4     featurewise_std_normalization=False, # divide inputs by std of the dataset  
5     samplewise_std_normalization=False, # divide each input by its std  
6     zca_whitening=False, # apply ZCA whitening  
7     rotation_range = 30, # randomly rotate images in the range (degrees, 0 to 180)  
8     zoom_range = 0.2, # Randomly zoom image  
9     width_shift_range=0.1, # randomly shift images horizontally (fraction of total width)  
10    height_shift_range=0.1, # randomly shift images vertically (fraction of total height)  
11    horizontal_flip = True, # randomly flip images  
12    vertical_flip=False) # randomly flip images  
13  
14  
15 datagen.fit(x_train)
```

- Define the Model

```
1 model = Sequential()
2 model.add(Conv2D(32,3,padding="same", activation="relu", input_shape=(224,224,3)))
3 model.add(MaxPool2D())
4
5 model.add(Conv2D(32, 3, padding="same", activation="relu"))
6 model.add(MaxPool2D())
7
8 model.add(Conv2D(64, 3, padding="same", activation="relu"))
9 model.add(MaxPool2D())
10 model.add(Dropout(0.4))
11
12 model.add(Flatten())
13 model.add(Dense(128,activation="relu"))
14 model.add(Dense(2, activation="softmax"))
15
16 model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 224, 224, 32)	896
max_pooling2d_1 (MaxPooling2D)	(None, 112, 112, 32)	0
conv2d_2 (Conv2D)	(None, 112, 112, 32)	9248
max_pooling2d_2 (MaxPooling2D)	(None, 56, 56, 32)	0
conv2d_3 (Conv2D)	(None, 56, 56, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 28, 28, 64)	0
dropout_1 (Dropout)	(None, 28, 28, 64)	0
flatten_1 (Flatten)	(None, 50176)	0
dense_1 (Dense)	(None, 128)	6422656
dense_2 (Dense)	(None, 2)	258

Total params: 6,451,554

Trainable params: 6,451,554

Non-trainable params: 0

- Model compilation and training

```
1 opt = Adam(lr=0.000001)
2 model.compile(optimizer = opt , loss = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True) ,
3               metrics = ['accuracy'])
```

```
1 history = model.fit(x_train,y_train,epochs = 500 , validation_data = (x_val, y_val))
```

ry libcublas.so.10

2022-11-05 15:26:33.581615: I tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully opened dynamic library libcudnn.so.7


2448/2448 [=====] - 11s 4ms/step - loss: 0.6933 - accuracy: 0.3877 - val_loss: 0.6936 - val_accuracy: 0.2547

Epoch 2/500

- Model evaluation after training.

```
1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3 loss = history.history['loss']
4 val_loss = history.history['val_loss']
5 epochs_range = range(500)
6
7 plt.figure(figsize=(15, 15))
8 plt.subplot(2, 2, 1)
9 plt.plot(epochs_range, acc, label='Training Accuracy')
10 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
11 plt.legend(loc='lower right')
12 plt.title('Training and Validation Accuracy')
13
14 plt.subplot(2, 2, 2)
15 plt.plot(epochs_range, loss, label='Training Loss')
16 plt.plot(epochs_range, val_loss, label='Validation Loss')
17 plt.legend(loc='upper right')
18 plt.title('Training and Validation Loss')
19 plt.show()
```

- Plot our training and validation accuracy along with training and validation loss



- Classification report.

```

i]: 1 predictions = model.predict_classes(x_val)
    2 predictions = predictions.reshape(1,-1)[0]
    3
    4 print(classification_report(y_val, predictions, target_names = ['Rugby (Class 0)','Soccer (Class 1)']))

```

	precision	recall	f1-score	support
Rugby (Class 0)	0.86	0.84	0.85	305
Soccer (Class 1)	0.73	0.76	0.74	170
accuracy			0.81	475
macro avg	0.80	0.80	0.80	475
weighted avg	0.81	0.81	0.81	475

```

16]: 1 model.evaluate(x_val,y_val)

```

610/610 [=====] - 1s 2ms/step

```

16]: [0.5176380370484024, 0.5622950792312622]

```

Transfer Learning

- Technique where a model trained on one task is re-purposed on a second related task.
- When the dataset is small, by using a pre-trained model helps to achieve high performance.
- Import the base model : VGG16
- Pre-trained on the ImageNet dataset, a large dataset consisting of 1.4M images and 1000 classes.
- This base of knowledge will help us classify Rugby and Soccer from our specific dataset.
- `include_top=False` argument : load a network that doesn't include the classification layers at the top.

```
1 base_model = tf.keras.applications.VGG16(input_shape = (224, 224, 3), include_top = False, weights = "imagenet")
2
```

```
[7]: 1 base_model = tf.keras.applications.VGG16(input_shape = (224, 224, 3), include_top = False, weights = "imagenet")
      2
```

```
[8]: 1 base_model.trainable = False
      2
```

```
[9]: 1 model = tf.keras.Sequential([base_model,
      2                             tf.keras.layers.GlobalAveragePooling2D(),
      3                             tf.keras.layers.Dropout(0.2),
      4                             tf.keras.layers.Dense(2, activation="softmax")
      5                             ])
```

```
[*]: 1 base_learning_rate = 0.00001
      2 model.compile(optimizer=tf.keras.optimizers.Adam(lr=base_learning_rate),
      3               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
      4               metrics=['accuracy'])
      5
      6 history = model.fit(x_train,y_train,epochs = 300 , validation_data = (x_val, y_val))
```

Train on 2448 samples, validate on 610 samples

Epoch 1/300

2448/2448 [=====] - 9s 4ms/sample - loss: 0.7029 - accuracy: 0.4726 - val_loss: 0.7071 - val_accuracy: 0.4148

Epoch 2/300

2448/2448 [=====] - 8s 3ms/sample - loss: 0.7020 - accuracy: 0.4808 - val_loss: 0.7046 - val_accuracy: 0.4180

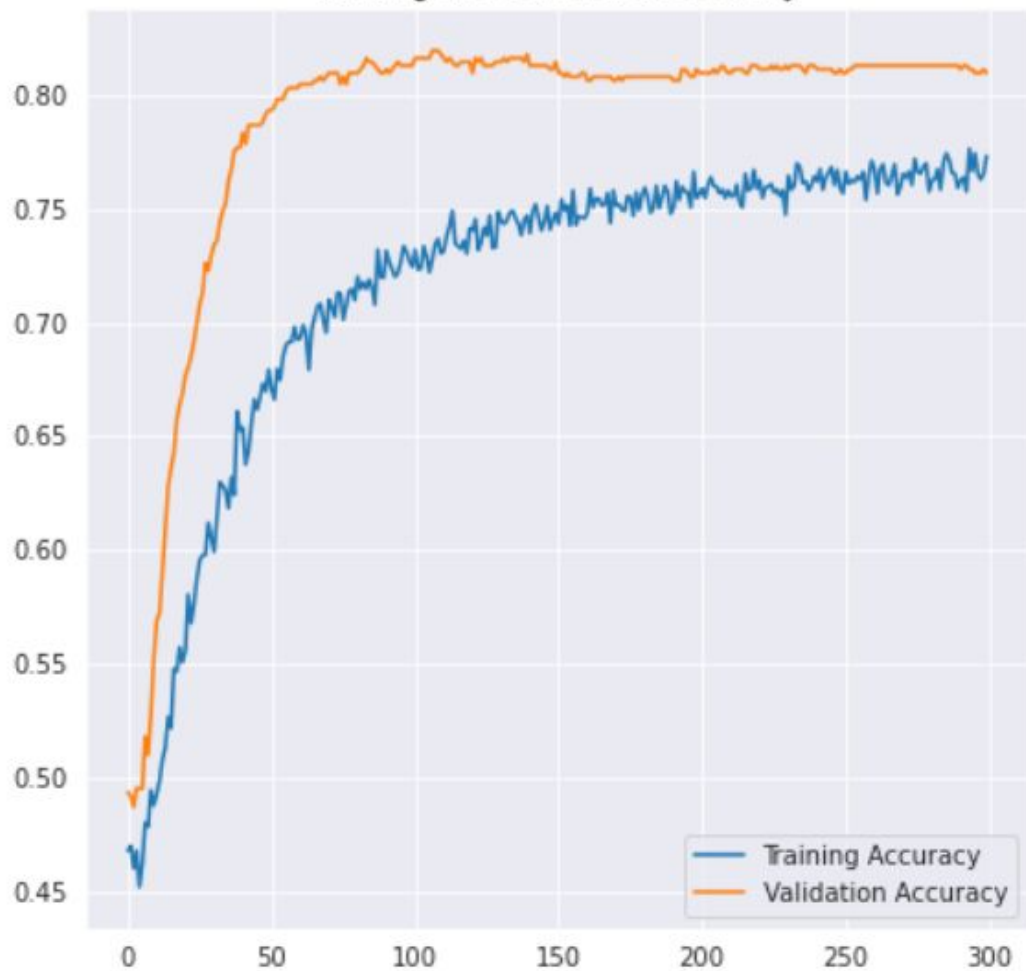
Epoch 3/300

- Model evaluation after training.

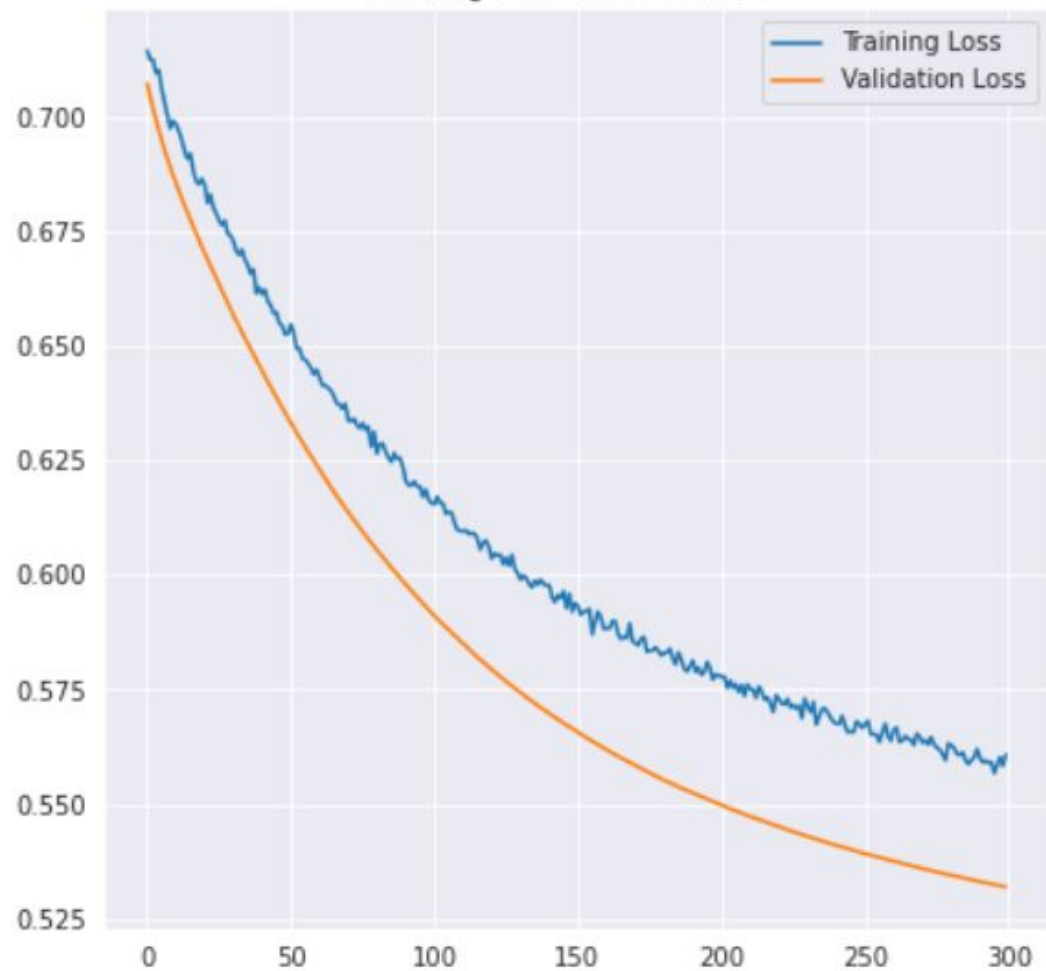
```
1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3 loss = history.history['loss']
4 val_loss = history.history['val_loss']
5 epochs_range = range(500)
6
7 plt.figure(figsize=(15, 15))
8 plt.subplot(2, 2, 1)
9 plt.plot(epochs_range, acc, label='Training Accuracy')
10 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
11 plt.legend(loc='lower right')
12 plt.title('Training and Validation Accuracy')
13
14 plt.subplot(2, 2, 2)
15 plt.plot(epochs_range, loss, label='Training Loss')
16 plt.plot(epochs_range, val_loss, label='Validation Loss')
17 plt.legend(loc='upper right')
18 plt.title('Training and Validation Loss')
19 plt.show()
```

- Plot our training and validation accuracy along with training and validation loss

Training and Validation Accuracy



Training and Validation Loss



- Classification report.

```
[1]: 1 predictions = model.predict_classes(x_val)
      2 predictions = predictions.reshape(1,-1)[0]
      3 print(classification_report(y_val, predictions, target_names = ['Rugby (Class 0)', 'Soccer (Class 1)']))
```

	precision	recall	f1-score	support
Rugby (Class 0)	0.79	0.89	0.83	305
Soccer (Class 1)	0.87	0.76	0.81	305
accuracy			0.82	610
macro avg	0.83	0.82	0.82	610
weighted avg	0.83	0.82	0.82	610

```
[16]: 1 model.evaluate(x_val,y_val)
```

```
610/610 [=====] - 1s 2ms/sample - loss: 0.5319 - accuracy: 0.8098
```

```
[16]: [0.5318700773794143, 0.8098361]
```


Thank You